Universitatea
Politehnica
București

Facultatea de
Automatică și
Calculatoare

Catedra de
Calculatoare

WRF ACCELERATION

AND COUPLING WITH FLEXPART

Author

Valentin Marcu – SCPD

ctvalentin.marcu@gmail.com

Supervisor

Associate Prof. Dr. Eng. Emil Slusanschi

Topics

- WRF Overview
- Radiation Transfers
- CUDA Overview
- FLEXPART Overview
- Related Work
- CUDA Acceleration
- Results
- Future Work
- Conclusions

WRF Overview

Weather Research and Forecasting Model

- Regional numeric model
- Scalability and efficiency
- MPI/OpenMP support
- NETCDF output

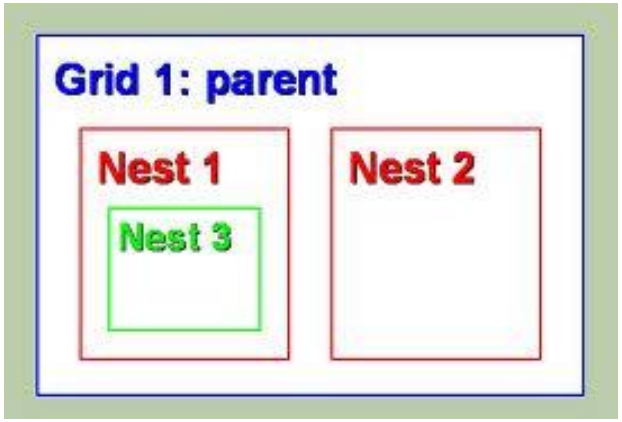

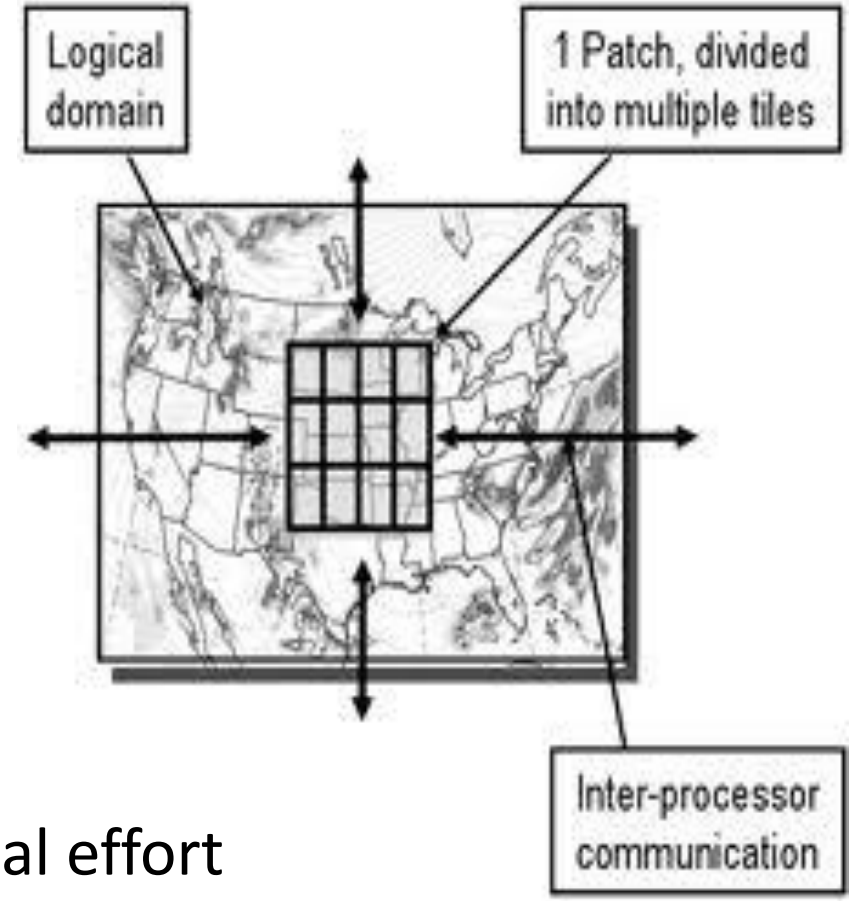

- Nest-based computational effort

WRF Overview

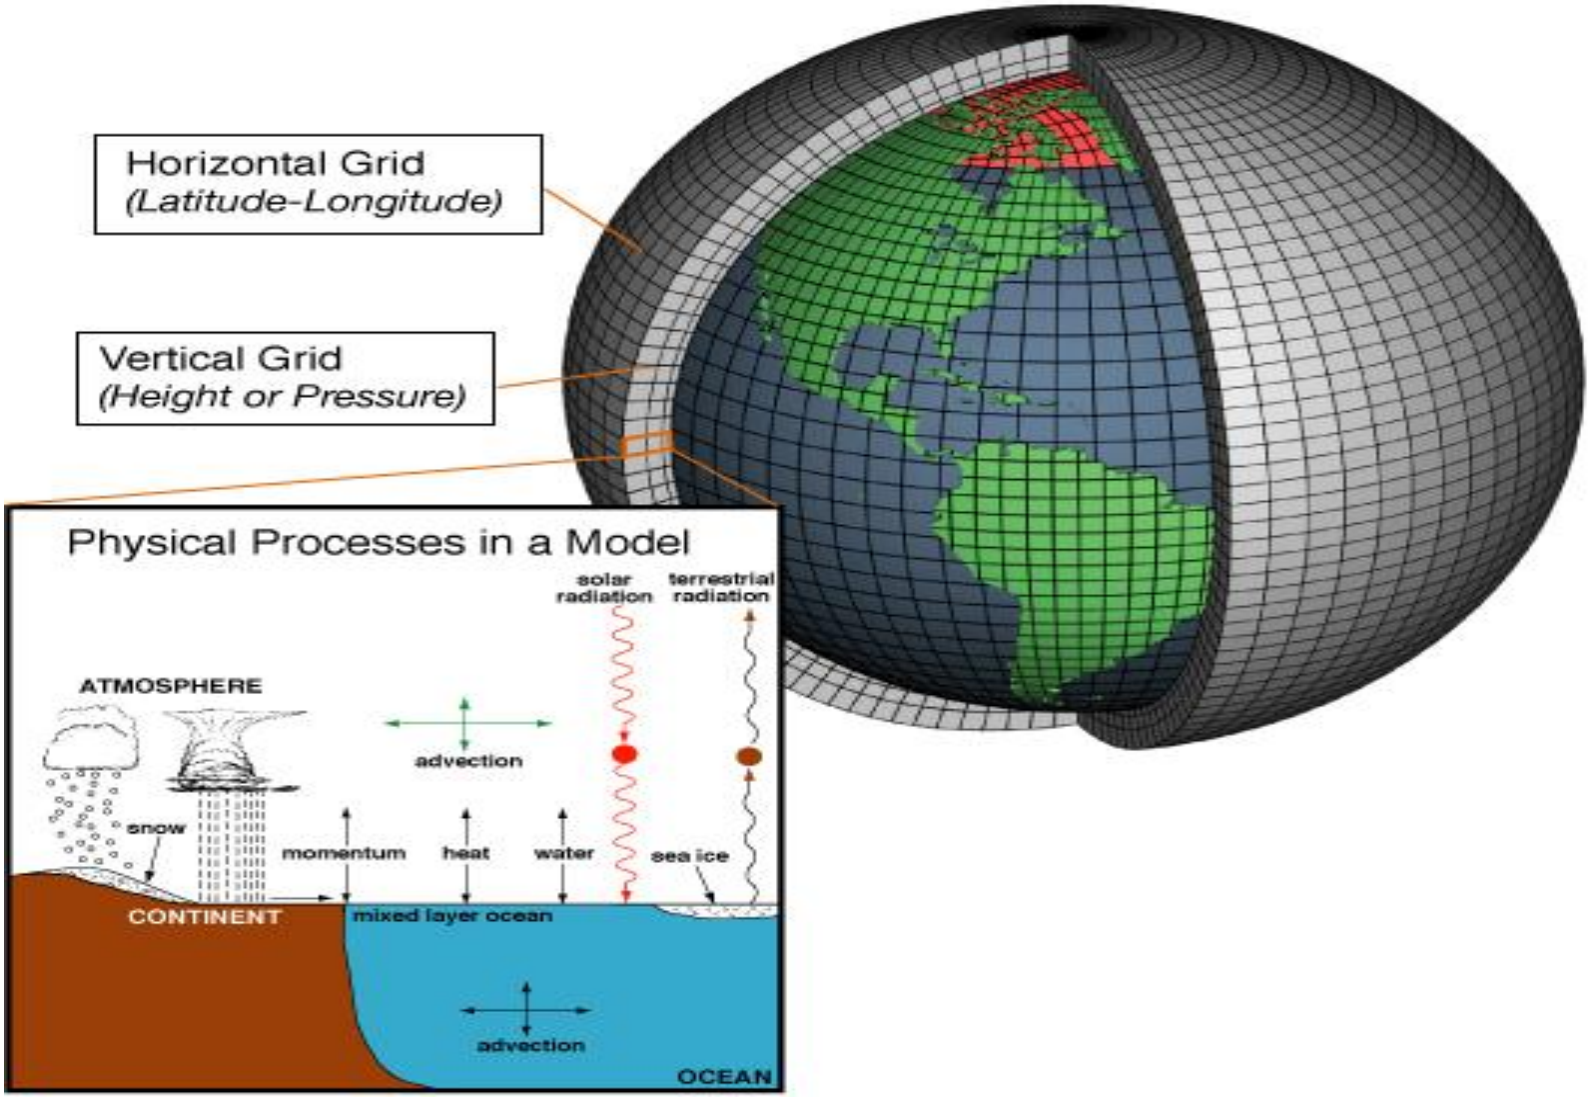

WRF Radiation Transfers

Rapid Radiative Transfer Model (RRTM)

- **inirad():** computes the ozone mixing ratio distribution
- **setcoef():** input for the radiative transfer algorithm
- **mm5atm():** prepares atmospheric profiles
- **gasabs():** calculates gaseous optical depths
  - taugb1()
  - taugb2 ()
  - …
  - taugb16 ()
- **rtrn():** calculates the radiative transfer

RRTM Call Graph

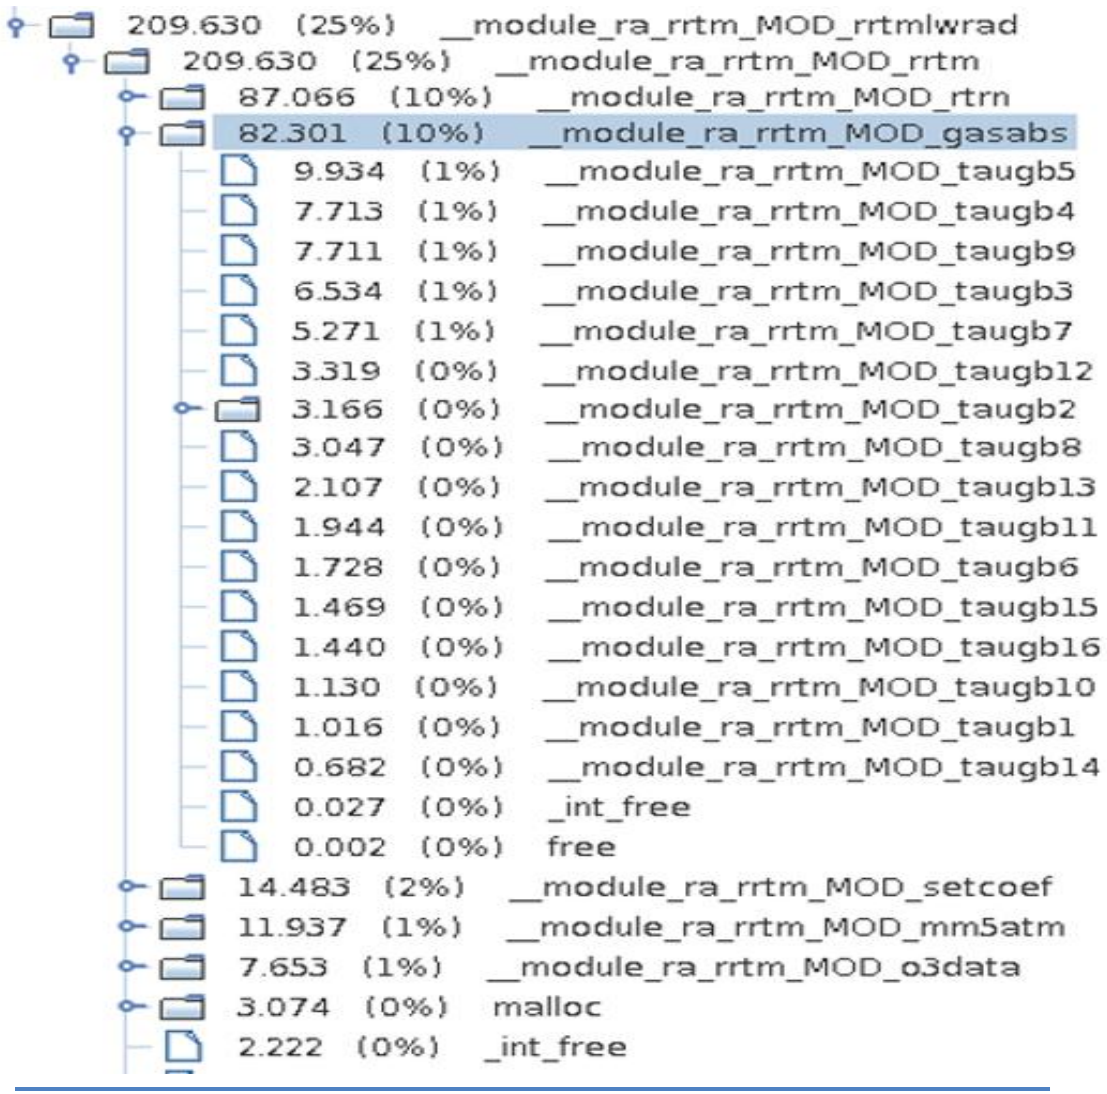

CUDA Overview

The GPU devotes more transistors to

data processing

(derived/inspired from graphics rendering)

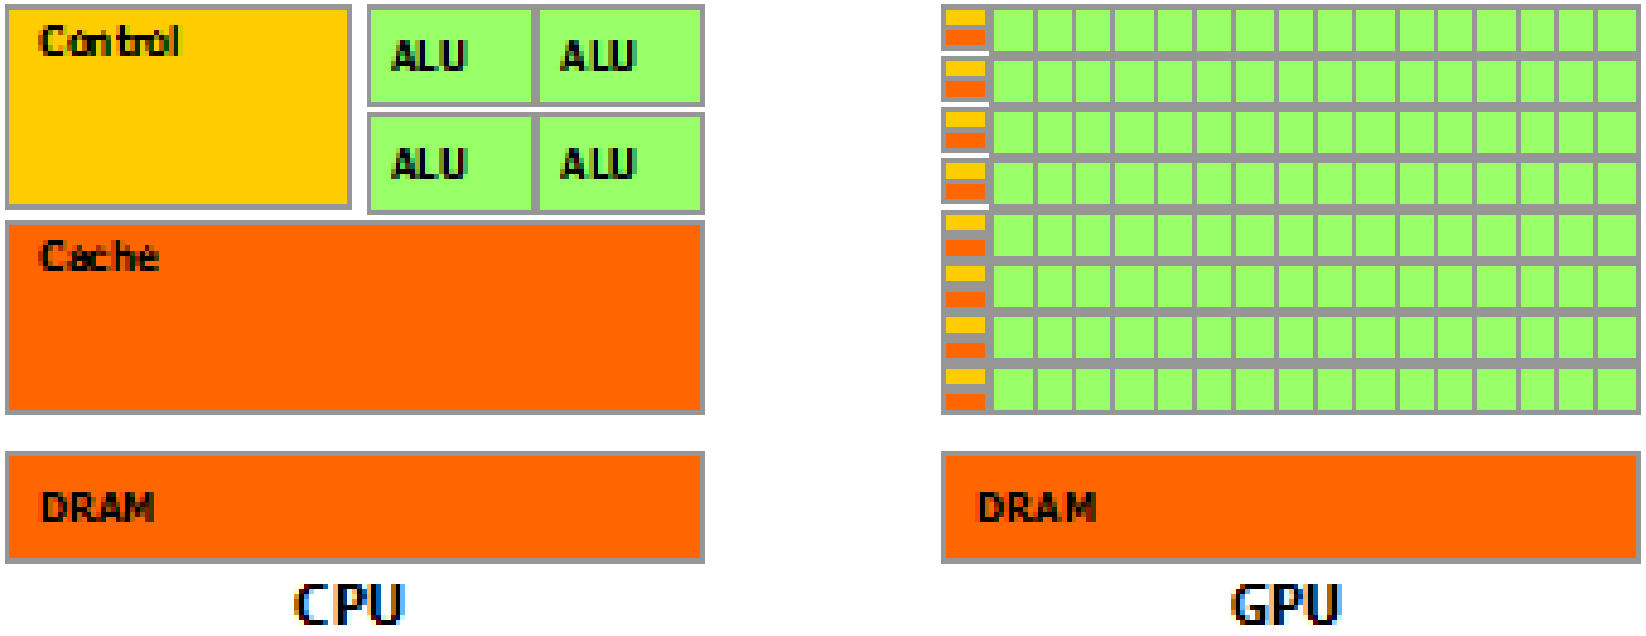

CUDA Memory Types

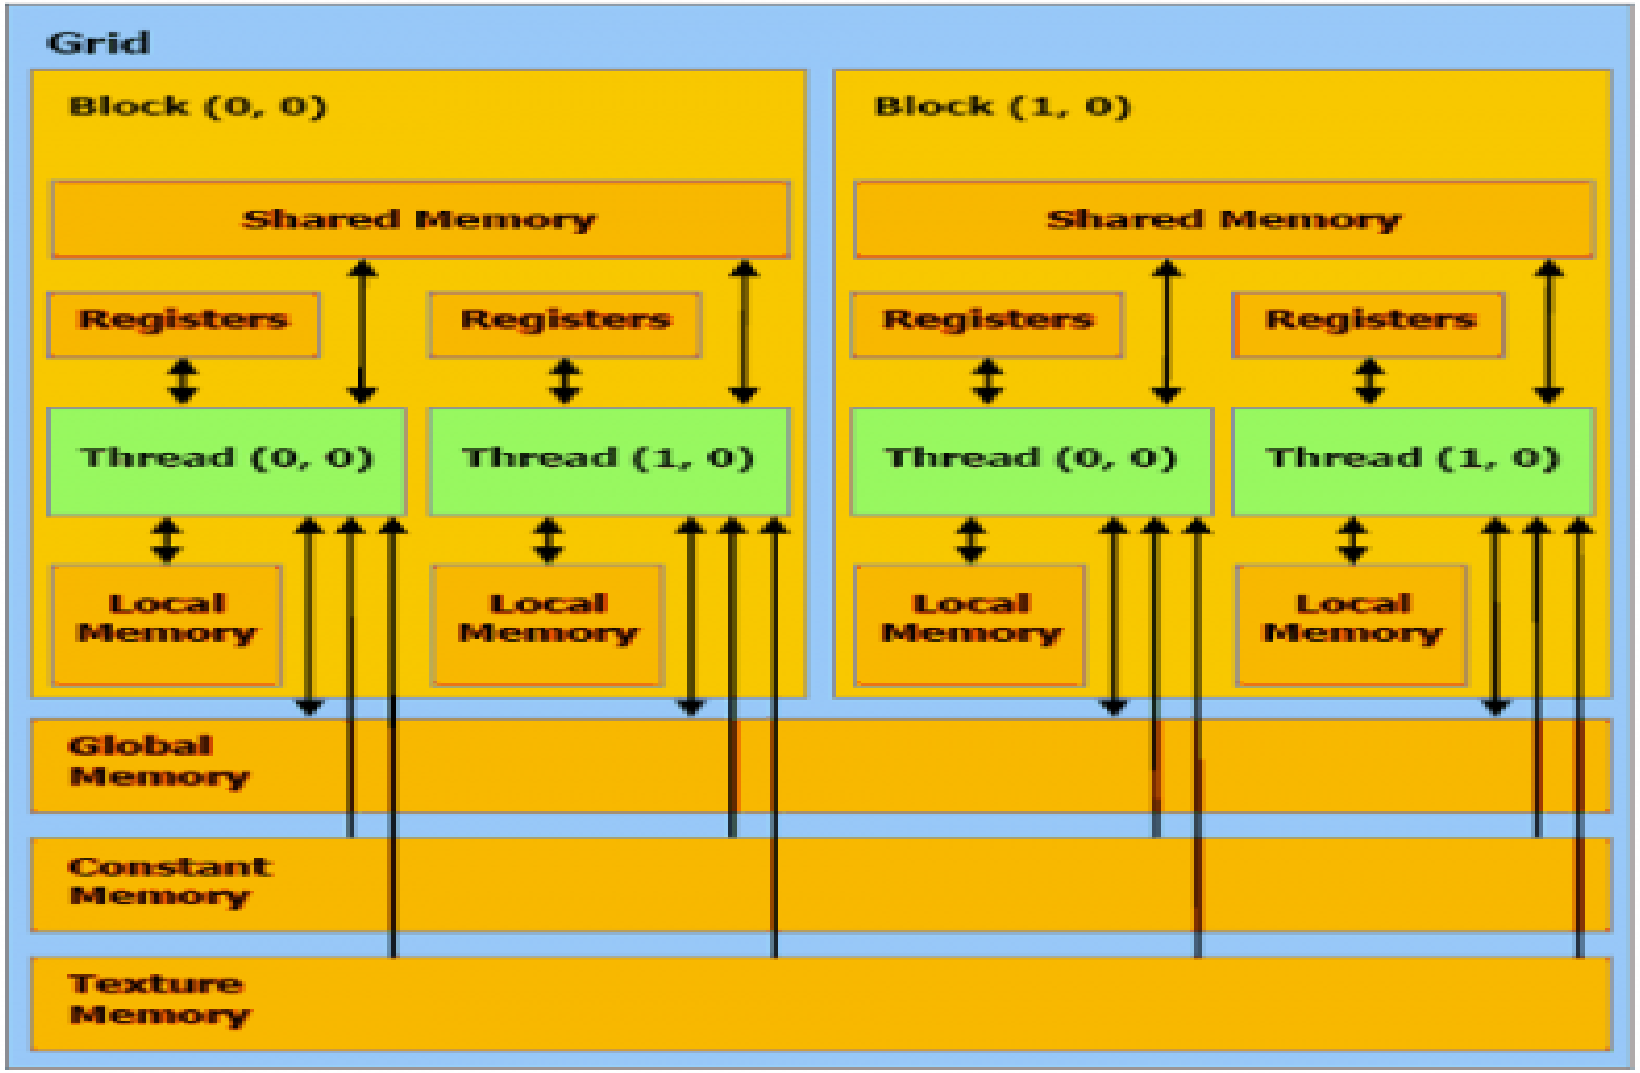

FLEXPART Overview

- Lagrangian particle dispersion model
- Originally designed for coupling with global weather models
- Forward/backward dispersion
- Radioactive decay, wet/dry deposition etc.
- Numerical parallelization for computation efficiency
- Netcdf format for output files

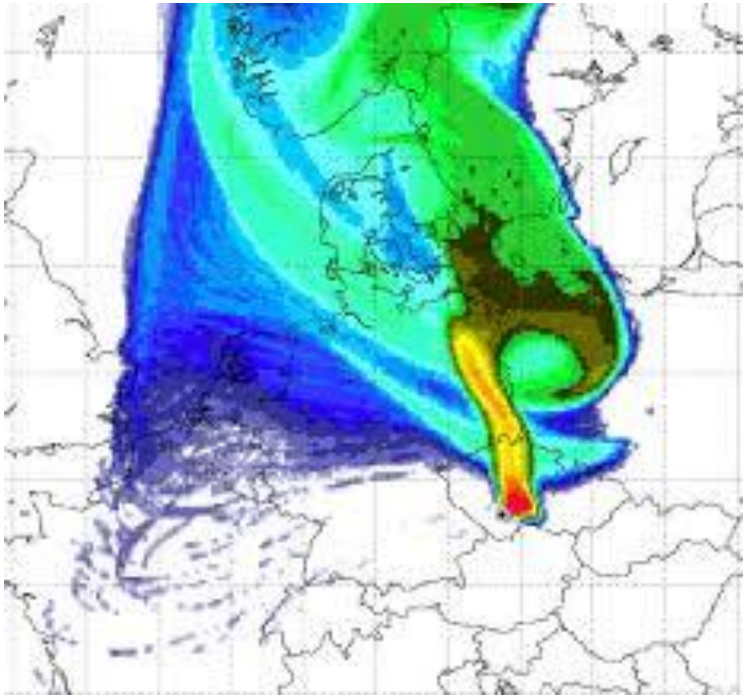

Related Work

GPU ACCELERATION OF RRTM USING CUDA FORTRAN

	X5540 (-O3)	X5540 (-fast)	C1060
Overall time (ms)	831	703	83

- G.Reutsch, NVIDIA Corporation, GPU acceleration of the Long-wave Rapid Radiative Transfer Model in WRF using CUDA Fortran, 2011
- WRF speedup: 1.25

CUDA Acceleration

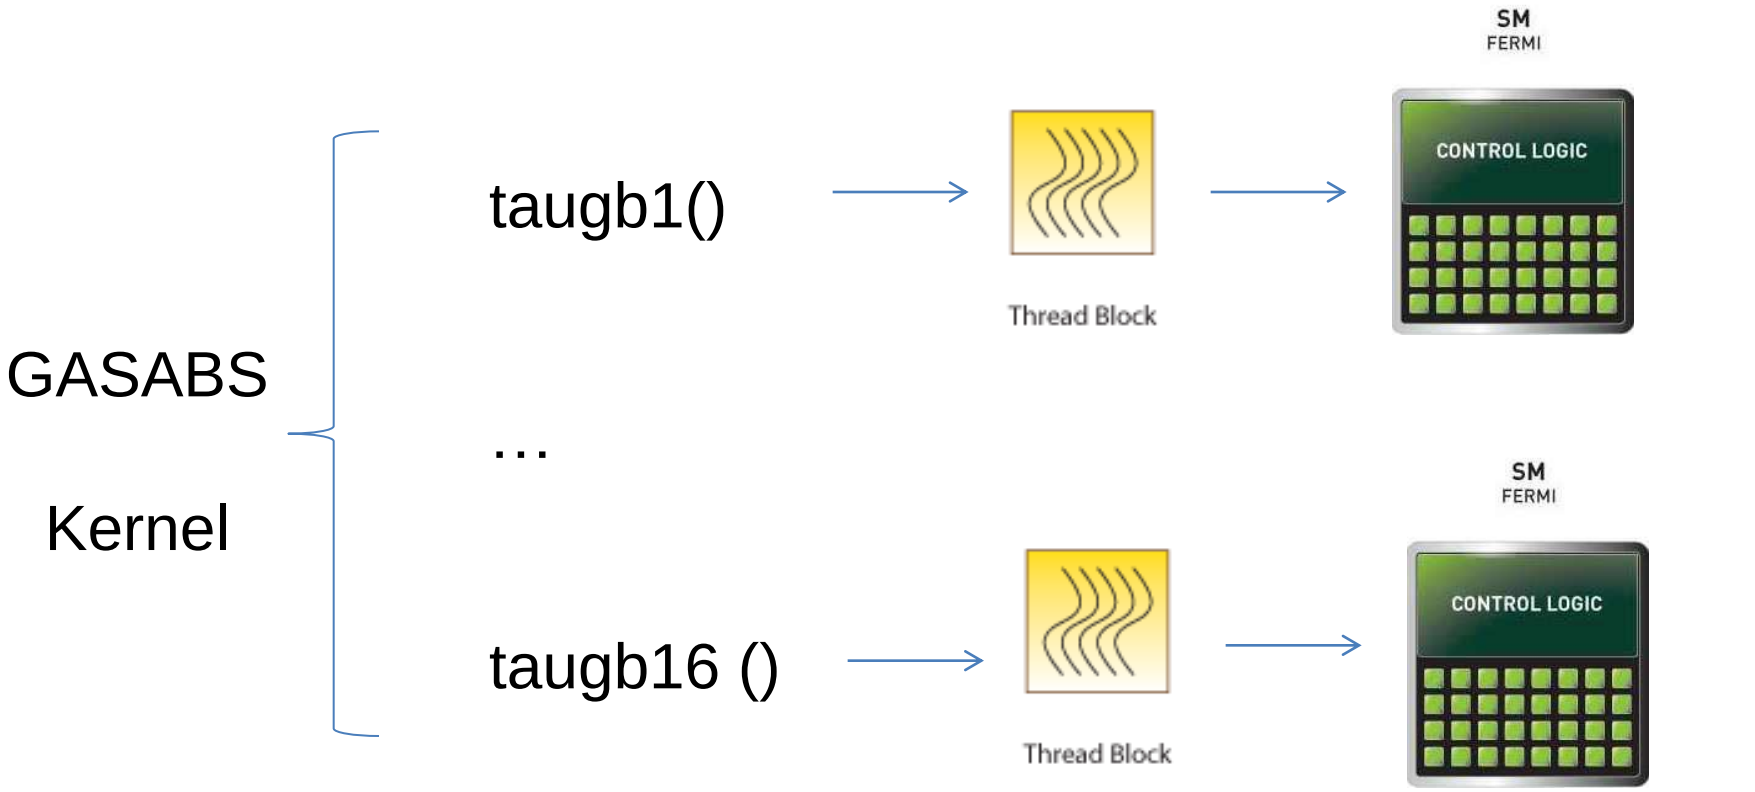

- 14 thread blocks
- 16 * NLAYERS threads / block
- 1 gasabs_gpu call for each 2D point in domain

- 14 SMs / device
- 2 devices / node

CUDA Acceleration

- 1 CUDA stream for each CPU thread (OpenMP)
- CUDA streams uniformly spread across GPUs
- 1 MPI process per node
- Asynchronous cudaMemcpy
- Overlap data transfers and kernel execution

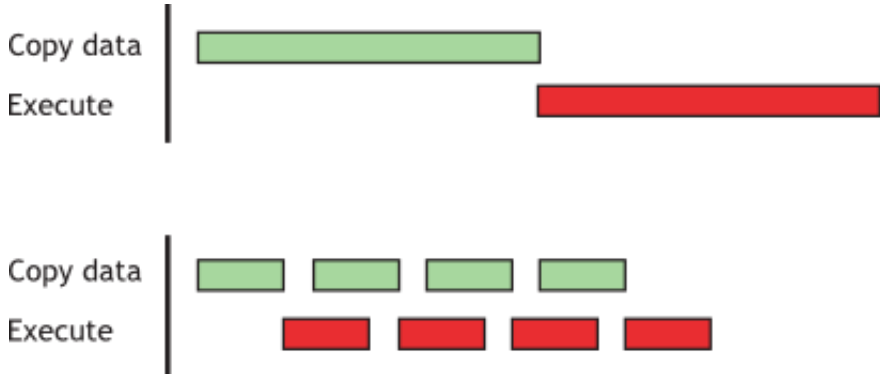

CUDA Acceleration

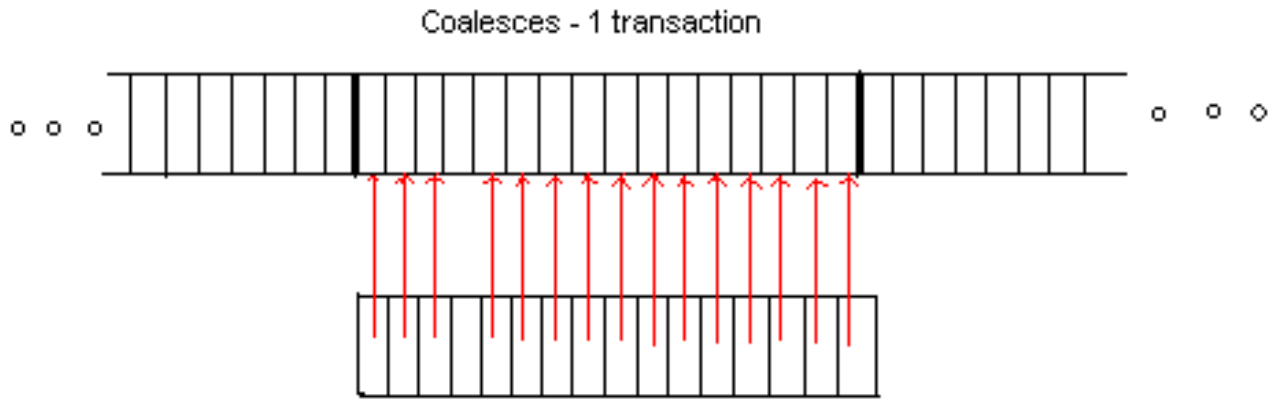

- Transposed some input arrays before sending them to GPU

  => coalesced accesses

  => minimized global memory transactions

- Allocated memory on GPU only at first kernel call

- Data transfers to/from GPU only when needed and in a minimum of transactions

Results (gfortran + nvcc)

- Europe Benchmark
- 1h of simulation
- 1 domain

Grid Points	Horiz Res
196x167x40 = 1.3M	21.6 km
274x217x40 = 2.4M	7.2 km
592x355x40 = 8.4M	2.4 km
1003x505x40 = 20.3M	800m

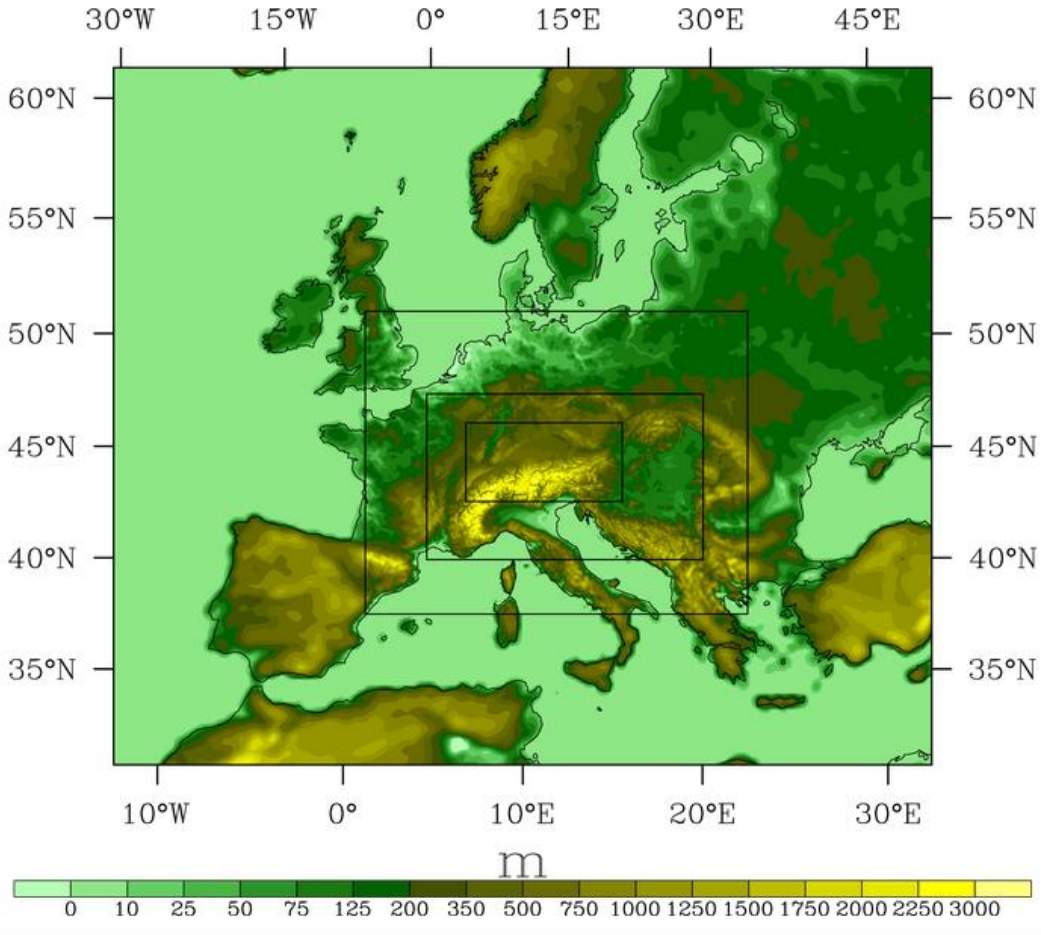

Results (gfortran + nvcc)

Percentage of Application

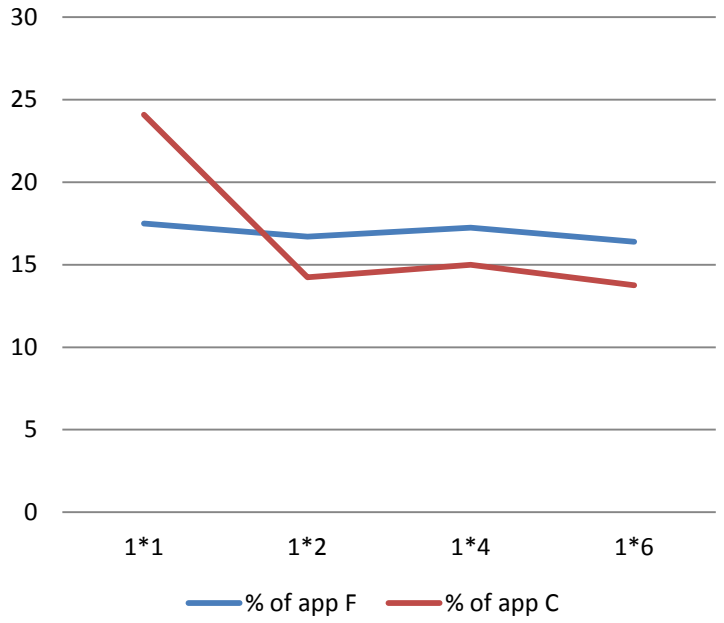

Total execution time

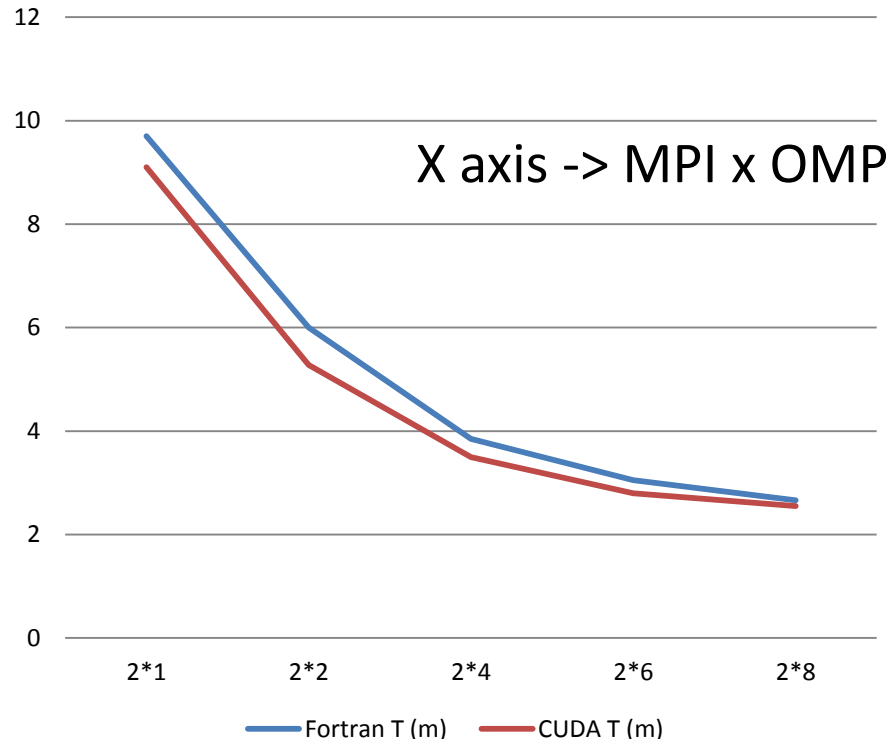


```
======== Profiling result:
 Time(%)      Time     Calls       Avg       Min       Max  Name
   53.26      1.03s     64740  15.96us  15.28us  17.88us  taugb_gpu(int, int,
   40.44  784.77ms     64740  12.12us  12.06us  13.92us  [CUDA memcpy DtoH]
    6.29  122.12ms     64752   1.89us   1.22us  92.80us  [CUDA memcpy HtoD]
```

FLEXPART – WRF Coupling

Domain Generation – WRF Domain Wizard

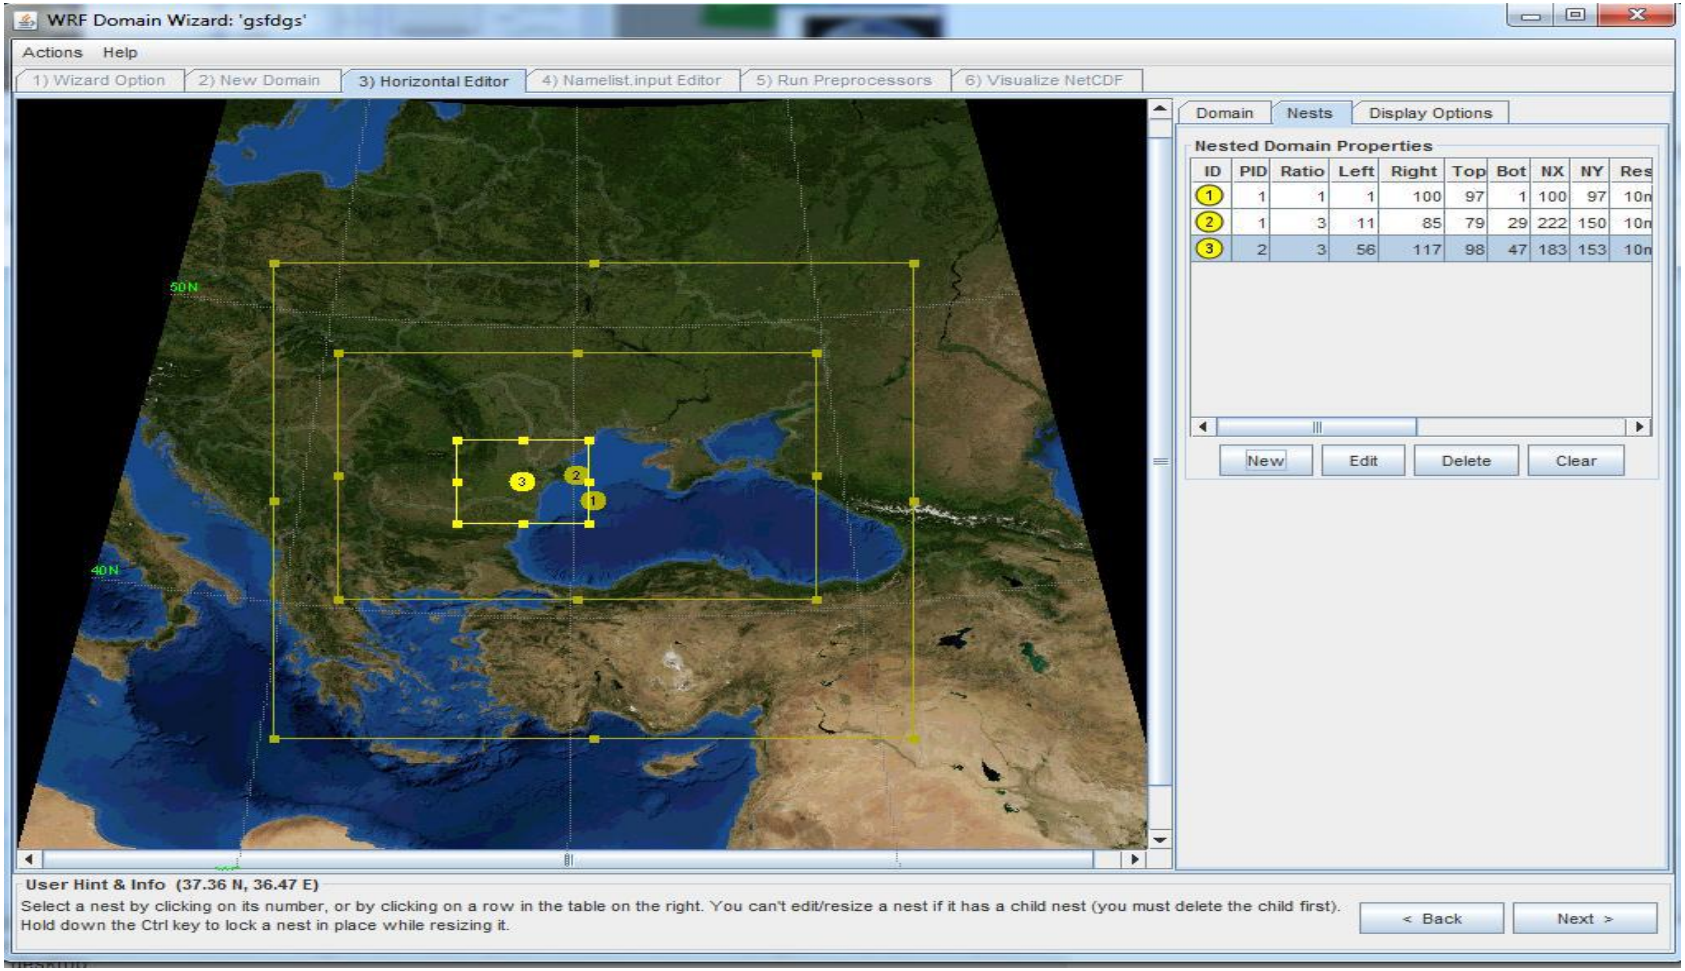

FLEXPART – WRF Coupling

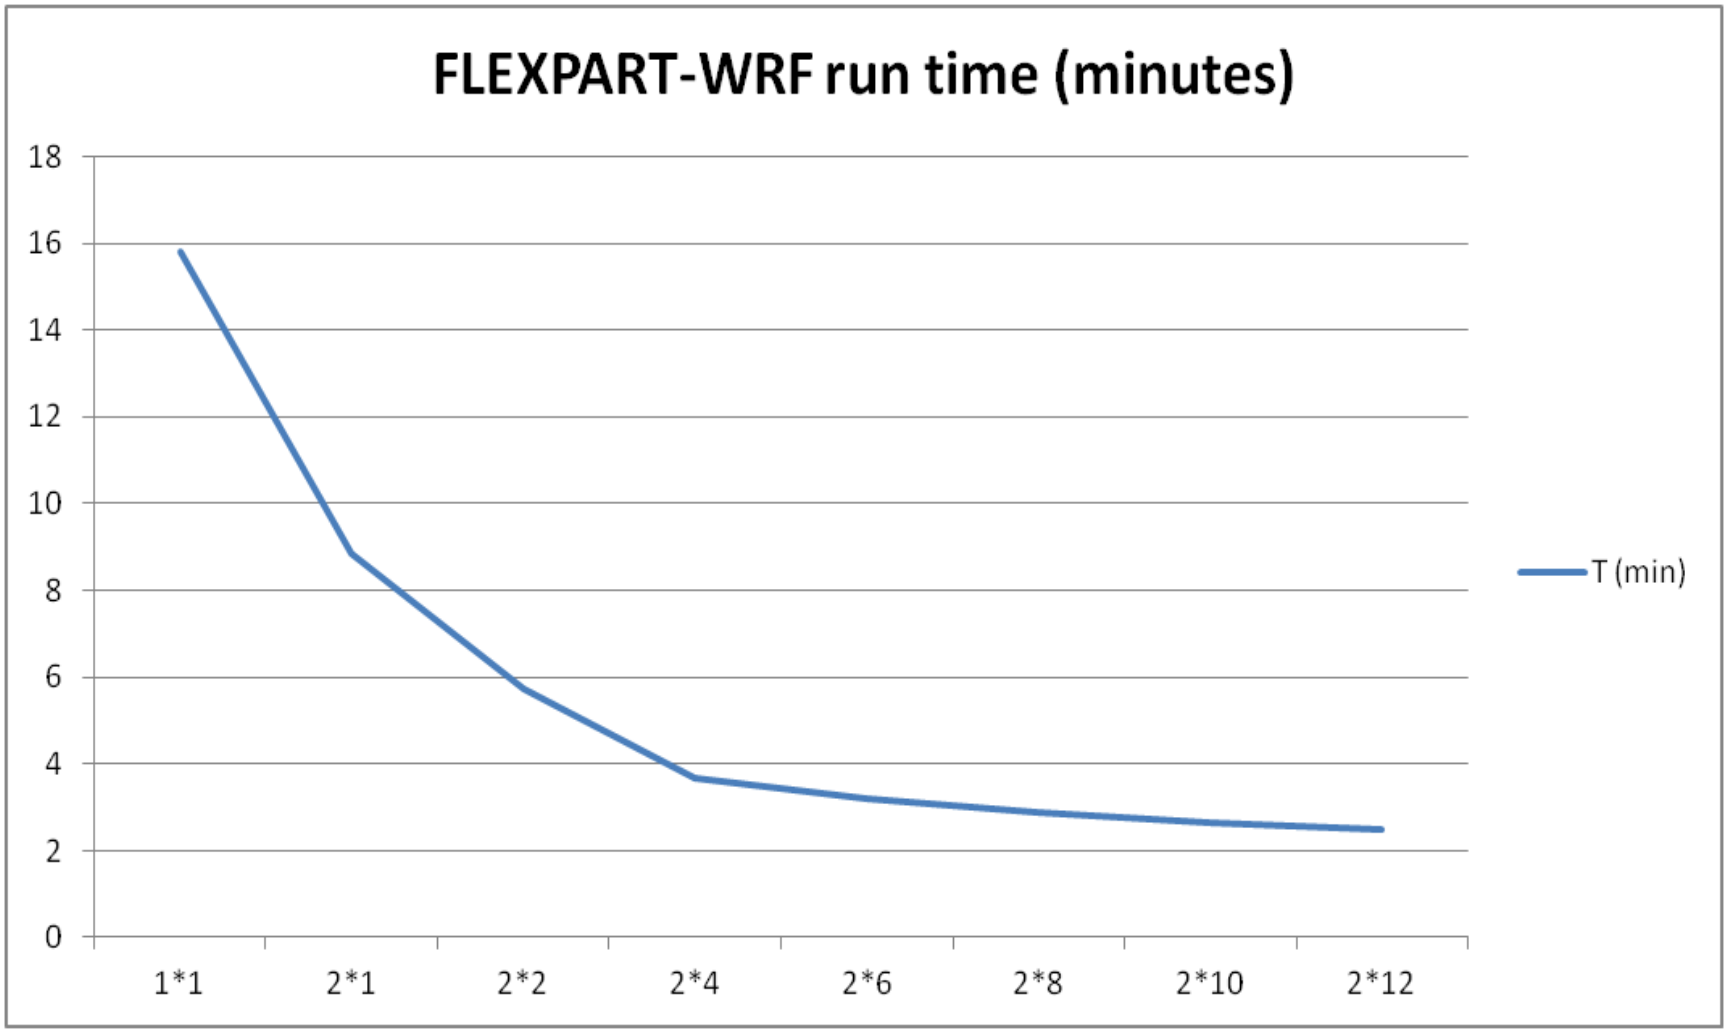

FLEXPART – WRF Coupling

Particle concentrations – coupled run

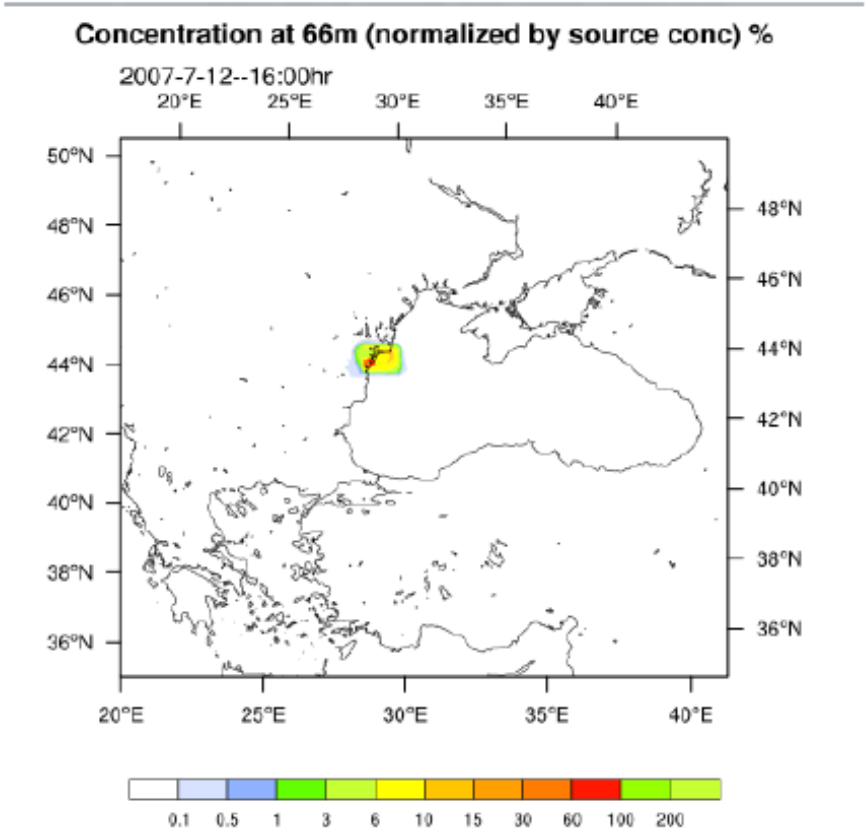

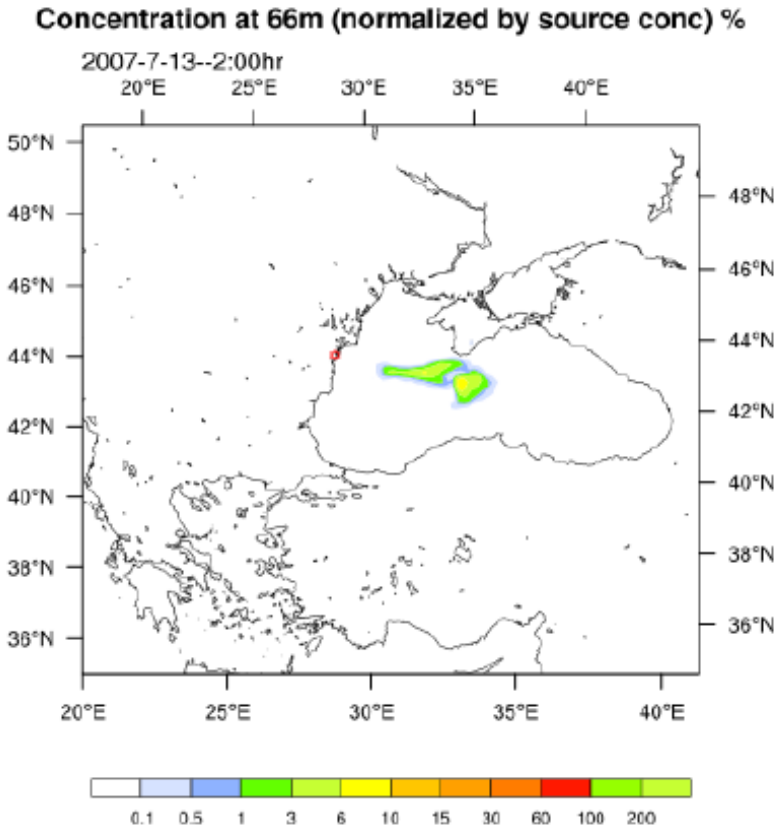

Future GPU Acceleration

OpenACC®

DIRECTIVES FOR ACCELERATORS

- Directive-based
- Multiple platforms supported (NVIDIA, AMD, Intel Xeon Phi, CRAY etc)
- v.2.0 (2014) – procedure calls

CPU-GPU transfers managed by the API

Conclusions

- Porting Fortran to CUDA C tends to result in better performance than using accelerator directives, but is very time consuming and prone to various implementation errors and vulnerabilities
- Overall WRF speedup of up to 1.1x
  - Gasabs() is only 10%-20% of WRF
- Upcoming accelerators and programming tools
- Scalability of last version of FLEXPART-WRF

Keywords

- WRF
- CUDA
- GPU
- RRTM
- Kernel
- Async
- FLEXPART
- Gfortran
- Domain
- OpenACC
- GASABS
- Speedup